

## high-level plan

### High-Level Project Plan: The Agentic App Genesis Engine

This plan outlines the architecture, phases, and technologies required to build a system that creates interactive web applications from natural language, inspired by v0.dev.

#### Project Vision

To create a multi-agent system that translates a user's abstract ideas into a functional web application prototype. The system will guide the user through a structured discovery process, visualize the application's architecture, and generate the initial codebase, with the ability to produce detailed development documentation for further expansion.

#### Core Architecture Philosophy

We will build this system on a **stateful, multi-agent architecture**. This is crucial because each step depends on the validated output of the previous one.

- **LangGraph** will be the backbone of our agentic flow. Its ability to create cyclical graphs allows agents to reflect, retry, and pass a shared "state" object (containing all user answers, wireframes, and decisions) between them. This is superior to a simple sequential chain for this complex task.
- **Pydantic** will be used extensively to define the structure of the data passed between agents. This ensures reliability and type safety for our state object, follow-up quizzes, and tool inputs.

---

### Development Phases & Agent Responsibilities

This project will be built in distinct phases, with each phase managed by a specialized agent (or a graph of agents).

#### Phase 1: The "Discovery & Scoping" Agent

This agent's primary goal is to transform a vague user request into a concrete set of application requirements. It executes the first three steps of your flow.

- **Task:**
  1. **Initial Input:** Present the user with two options: a free-form text input or a guided, branching quiz to select a general app category (e.g., "e-commerce," "social media," "productivity tool").
  2. **Clarification & Enrichment:**
    - It analyzes the initial input. An internal "prompt evaluator" checks for key information.
    - If the input is insufficient, the agent asks up to three targeted follow-up questions. **These questions will be dynamically generated based on a master prompt.**
  3. **Structured Confirmation:**
    - Based on the answers, the agent generates two rounds of structured, multiple-choice quizzes (using Pydantic models to define the format) to lock in key decisions (e.g., "Target Audience: A) General Public, B) Businesses", "Monetization: A) Free, B) Subscription").
- **System Instruction Snippet for this Agent's LLM:**

You are a friendly and insightful Product Manager AI. Your goal is to help a user define their app idea.

- If the user's request is vague, you MUST ask clarifying questions about the app's PURPOSE, DESIGN STYLE, and KEY FEATURES.
- You will then generate structured multiple-choice questions to confirm these details.
- Your final output must be a JSON object containing all the user's validated choices.

- **Output:** A rich JSON object ( `app_brief.json` ) containing a structured summary of the user's requirements. This object becomes part of the main state passed to the next agent.

## Phase 2: The "UI/UX Architect" Agent

This agent takes the confirmed requirements and visualizes the application's structure and flow.

- **Task:**
  1. Ingest the `app_brief.json` from the previous phase.
  2. Generate two distinct versions of **wireframes** for the main screens (e.g., homepage, dashboard).
  3. Generate two distinct versions of **user flow diagrams** for the core user journey (e.g., sign-up process, posting content).
  4. Present these visual options to the user for selection.
- **Technology for Image Generation:**
  - To fulfill this, you should use a powerful multimodal model. **Google's Gemini 1.5 Pro** is perfectly suited for this. It can understand the complex request and generate visual representations.
  - **Practical Implementation:** Instead of direct pixel-perfect images, it's more robust to have Gemini output a description in a format that can be rendered, such as:
    - **Mermaid.js or PlantUML syntax:** This is text-based code that can be easily rendered into clean flowcharts and diagrams.
    - **SVG code:** A vector format that can be directly embedded in HTML.
    - **A detailed JSON structure** describing the layout, which a frontend component can then render.
- **System Instruction Snippet:**

You are a Senior UI/UX Designer AI. Your input is a JSON app brief.

- Your task is to create wireframes and user flow diagrams based on this brief.
- You MUST output your diagrams in Mermaid.js syntax.
- For wireframes, you MUST output a structured JSON describing the layout of UI components (e.g., `{ "type": "header", "children": [...] }, { "type": "button", "label": "Sign Up" }`).
- Provide two distinct versions for each deliverable.

- **Output:** The selected wireframe/flow versions are added to the state object.

## Phase 3: The "Code Generation" Agent

This is the final execution agent for the MVP. It synthesizes all prior decisions into a tangible code prototype.

- **Task:**
  1. Receive the final state object containing the app brief and the chosen UI/UX designs.
  2. Translate the structured wireframe JSON and user flows into code.
  3. Generate the code for **HTML, CSS, and JavaScript** files. The AI should create a clean, modular structure (e.g., `index.html`, `style.css`, `app.js`).
  4. Pass the generated file structure and content to the Sandbox for rendering.
- **System Instruction Snippet:**

You are an Expert Frontend Developer AI. Your input is a final project specification including a JSON-based wireframe definition.

- You must write clean, standards-compliant HTML, CSS, and vanilla JavaScript.
- Create separate files for structure (HTML), style (CSS), and logic (JS).
- Your output MUST be a JSON object where keys are filenames (e.g., "index.html") and values are the code content as strings.
- Do not invent new features; strictly adhere to the provided specification.

- **Output:** A JSON object representing the file system of the generated code.

## Phase 4: The "Documentation & Roadmap" Agent (Post-MVP)

This agent provides the high-level documentation needed to turn the prototype into a real project.

- **Task:**
  1. Based on the final project state, generate a series of professional documents.
  2. The user can select which documents to create from a list.
- **Deliverables:**
  - **Higher-Level Plan:** A strategic overview.
  - **Project Architecture:** A technical diagram and explanation.
  - **Product Requirements Document (PRD):** A formal doc for product teams.
  - **Tech Specification:** Detailed spec for developers.
  - **Project Phases & Roadmap:** A timeline for development.
  - **Development Sprints & Tasks:** Break down the project into actionable sprints and tasks with sample code snippets.

## Technology Stack & Key Decisions

- **Backend & Agent Orchestration:**
  - **Framework: FastAPI** (Python). It's modern, fast, and integrates seamlessly with the Python AI ecosystem.
  - **Agent Logic: LangGraph** to manage the stateful, multi-step flow between agents.
  - **Data Validation: Pydantic** to define all data structures.
- **AI & LLMs:**
  - **Core Logic & Text: Claude 3.5 Sonnet** or **GPT-4o** are excellent choices for reasoning and structured data generation.
  - **Multimodality (Phase 2): Gemini 1.5 Pro** for generating the wireframes and user flows.
- **Sandbox for Rendering:**
  - **The Problem:** You need to safely execute the generated code (HTML/CSS/JS) and display the result.
  - **Recommended Solution (for MVP): Docker.** It's the most straightforward and robust way to start. Your FastAPI backend can use the `docker-py` library to:
    1. Create a temporary directory with the generated files.
    2. Create a simple `Dockerfile` on the fly to serve these static files (e.g., using a basic Python HTTP server or Nginx).
    3. Build a Docker image and run a container from it, mapping a port.
    4. The frontend can then display the result from that container's URL in an `<iframe>`.
- **Frontend (For the Agentic App itself):**
  - **Framework: Next.js (React)** with TypeScript.
  - **UI Kit: Shadcn/UI** and **Tailwind CSS** for a modern, component-based interface.